



Introduction to **Machine Learning and Data Mining** (Học máy và Khai phá dữ liệu)

Khoat Than

School of Information and Communication Technology
Hanoi University of Science and Technology


2020

Contents


- Introduction to Machine Learning & Data Mining
- **Supervised learning**
 - **K-nearest neighbors**
- Unsupervised learning
- Practical advice

Which class does the object belong to?


Class a




Class a




Class b




Class a




?



Class a



Class b



Neighbor-based learning (1)


- *K-nearest neighbors* (KNN) is one of the most simple methods in ML. Some other names:
 - Instance-based learning
 - Lazy learning
 - Memory-based learning
- *Main ideas:*
 - There is no specific assumption on the function to be learned.
 - Learning phase just stores all the training data.
 - Prediction for a new instance is based on *its nearest neighbors* in the training data.
- Thus KNN is called a *non-parametric method*.
(no specific assumption on the classifier/regressor)

Neighbor-based learning (2)

- Two main ingredients:
 - The **similarity measure** (distance) between instances/objects.
 - The **neighbors** to be taken in prediction.
- *Under some conditions, KNN can achieve the Bayes-optimal error which is the performance limit of any methods.* [Györfi and Hengartner, JMLR 2013]
 - Even 1-NN (with some simple modifications) can reach this performance. [Kontorovich & Weiss, AISTATS 2015]
- KNN is close to **Manifold learning**.

KNN: example

- Take 1 nearest neighbor?
 - Assign \mathbf{z} to class 2.
- Take 3 nearest neighbors?
 - Assign \mathbf{z} to class 1.
- Take 5 nearest neighbors?
 - Assign \mathbf{z} to class 1.



KNN for classification

- Data representation:
 - Each observation is represented by a vector in an n -dimensional space, e.g., $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{in})^T$. Each dimension represents an attribute/feature/variable.
 - There is a set C of predefined labels.
- Learning phase:
 - Simply save all the training data \mathbf{D} , with their labels.
- Prediction: for a new instance \mathbf{z} .
 - *For each instance \mathbf{x} in \mathbf{D} , compute the distance/similarity between \mathbf{x} and \mathbf{z} .*
 - *Determine a set $NB(\mathbf{z})$ of the nearest neighbors of \mathbf{z} .*
 - *Using majority of the labels in $NB(\mathbf{z})$ to predict the label for \mathbf{z} .*

KNN for regression

■ Data representation:

- Each observation is represented by a vector in an n -dimensional space, e.g., $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{in})^T$. Each dimension represents an attribute/feature/variable.
- The output y is a real number.

■ Learning phase:

- Simply save all the training data \mathbf{D} , with their labels.

■ Prediction: for a new instance \mathbf{z} .

- For each instance \mathbf{x} in \mathbf{D} , compute the distance/similarity between \mathbf{x} and \mathbf{z} .
- Determine a set $NB(\mathbf{z})$ of the nearest neighbors of \mathbf{z} , with $|NB(\mathbf{z})| = k$.
- Predict the label for \mathbf{z} by $y_z = \frac{1}{k} \sum_{x \in NB(z)} y_x$

KNN: two key ingredients (1)




Different
thoughts,

Different
views

Different
measures


KNN: two key ingredients (2)

- The **distance/similarity** measure
 - Each measure implies a view on data
 - Infinite many measures !!!
 - What measure should be?



KNN: two key ingredients (3)

- The set $NB(\mathbf{z})$ of **nearest neighbors**.
 - How many neighbors are enough?
 - How can we select $NB(\mathbf{z})$?
(by choosing k or restricting the area?)



KNN: 1 or more neighbors?

- In theory, *1-NN can be among the best methods under some conditions*. [Kontorovich & Weiss, AISTATS 2015]
- KNN is Bayes optimal under some conditions: Y bounded, large training size M , and the true regression function being continuous, and

$$k \rightarrow \infty, (k/M) \rightarrow 0, (k/\log M) \rightarrow +\infty$$

- In practice, we should use more neighbors for prediction ($k > 1$), but not too many:
 - To avoid noises/errors in only one nearest neighbor.
 - Too many neighbors might break the inherent structure of the data manifold, and thus prediction might be bad.

Distance/similarity measure (1)

- The distance measure:
 - Plays a very important role in KNN.
 - Indicates **how we assume/suppose the distribution of our data.**
 - Be determined once, and does not change in all prediction later.
- Some common distance measures:
 - *Geometric distance*: usable for problems with real inputs
 - *Hamming distance*: usable for problems with binary inputs, such as x in $\{0; 1\}$

Distance/similarity measure (2)

- Some geometric distances:

- Minkowski (L_p -norm):

$$d(x, z) = \left(\sum_{i=1}^n |x_i - z_i|^p \right)^{1/p}$$

- Manhattan (L_1 -norm):

$$d(x, z) = \sum_{i=1}^n |x_i - z_i|$$

- Euclid (L_2 -norm):

$$d(x, z) = \sqrt{\sum_{i=1}^n (x_i - z_i)^2}$$

- Chebyshev (max norm):

$$\begin{aligned} d(x, z) &= \lim_{p \rightarrow \infty} \left(\sum_{i=1}^n |x_i - z_i|^p \right)^{1/p} \\ &= \max_i |x_i - z_i| \end{aligned}$$

Distance/similarity measure (3)

- Hamming distance:

- for problems with binary inputs
- such as $\mathbf{x} = (1,0,0,1,1)$

$$d(x, z) = \sum_{i=1}^n \text{Difference}(x_i, z_i)$$

$$\text{Difference}(a, b) = \begin{cases} 1, & \text{if } (a \neq b) \\ 0, & \text{if } (a = b) \end{cases}$$

- Cosine measure:

- Suitable for some problems with textual inputs.

$$d(x, z) = \frac{x^T z}{\|x\| \cdot \|z\|}$$

KNN: attribute normalization

- Normalizing the attributes is sometimes important to get good predictiveness in KNN.
 - No normalization implies that the magnitude of an attribute might play a heavy role, and artificially overwhelms the other attributes. Ex.:

$$d(x, z) = \sqrt{\sum_{i=1}^n (x_i - z_i)^2}$$

- $x = (\text{Age}=20, \text{Income}=12000, \text{Height}=1.68)$
- $z = (\text{Age}=40, \text{Income}=1300, \text{Height}=1.75)$
- $d(x, z) = [(20-40)^2 + (12000-1300)^2 + (1.68-1.75)^2]^{0.5}$

- This is unrealistic and unexpected in some applications.
- Some common normalizations:
 - Make all values of x_j in $[-1; 1]$;
 - Make all values of x_j to have empirical mean 0 and variance 1.

KNN: attribute weighting

- Weighting the attributes is sometimes important for KNN.
 - No weight implies that the attributes play an equal role, e.g., due to the use of the Euclidean distance:


$$d(x, z) = \sqrt{\sum_{i=1}^n (x_i - z_i)^2} \quad \longrightarrow \quad d(x, z) = \sqrt{\sum_{i=1}^n w_i (x_i - z_i)^2}$$

- This is unrealistic in some applications, where an attribute might be more important than the others in prediction.
 - Some weights (w_i) on the attributes might be more suitable.
- How to decide the weights?
 - Base on the knowledge domain about your problem.
 - Learn the weights automatically from the training data.

KNN: weighting neighbors (1)

$$y_z = \frac{1}{k} \sum_{x \in NB(z)} y_x$$

- Prediction of labels miss some information about neighbors.
 - The neighbors in $NB(\mathbf{z})$ play the same role with respect to the different distances to the new instance.
 - This is unrealistic in some applications, where *closer neighbors should play more important role than the others*.
- Using the distance as weights in prediction might help.
 - Closer neighbors should have more effects.
 - Farther points should have less effects.



KNN: weighting neighbors (2)

- Let v be the weights to be used.

- $v(\mathbf{x}, \mathbf{z})$ can be chosen as the inverse of the distance from \mathbf{x} to \mathbf{z} , $d(\mathbf{x}, \mathbf{z})$.

- Some examples:

$$v(x, z) = \frac{1}{\alpha + d(x, z)} \quad v(x, z) = \frac{1}{\alpha + [d(x, z)]^2} \quad v(x, z) = e^{-\frac{d(x, z)^2}{\sigma^2}}$$

- For classification:

$$c_z = \arg \max_{c_j \in C} \sum_{x \in NB(z)} v(x, z) \cdot \text{Identical}(c_j, c_x)$$

$$\text{Identical}(a, b) = \begin{cases} 1, & \text{if } (a = b) \\ 0, & \text{if } (a \neq b) \end{cases}$$

- For regression:

$$y_z = \frac{\sum_{x \in NB(z)} v(x, z) \cdot y_x}{\sum_{x \in NB(z)} v(x, z)}$$

KNN: limitations/advantages

■ Advantages:

- *Low cost for the training phase.*
- *Very flexible in choosing the distance/similarity measure: we can use many other measures, such as Kullback-Leibler divergence, Bregman divergence,...*
- *KNN is able to reduce some bad effects from noises when $k > 1$.*
- *In theory, KNN can reach the best performance among all regression methods, under some conditions.
(this might not be true for other methods)*

■ Limitations:

- Have to find a suitable distance/similarity measure for your problem.
- Prediction requires intensive computation.

References

- A. Kontorovich and Weiss. A Bayes consistent 1-NN classifier. *Proceedings of the 18th International Conference on Artificial Intelligence and Statistics (AISTATS)*. JMLR: W&CP volume 38, 2015.
- A. Guyader, N. Hengartner. On the Mutual Nearest Neighbors Estimate in Regression. *Journal of Machine Learning Research* 14 (2013) 2361-2376.
- L. Gottlieb, A. Kontorovich, and P. Nisnevitch. Near-optimal sample compression for nearest neighbors. *Advances in Neural Information Processing Systems*, 2014.

Exercises

- What is the different between KNN and OLS?
- Is KNN prone to overfitting?
- How to make KNN work with sequence data?
(each instance is a sequence)